

Parallelisation of Graph Algorithms

in Julia

Speaking Order

1. Luke Thompson
2. Darcy Cox
3. George Qiao

Group 5

OVERVIEW



PARALLEL COMPUTING IN JULIA



PARALLELISATION OF GRAPH ALGORITHMS




PARALLEL GRAPH IMPLEMENTATION & BENCHMARKING


Julia Language

What is Julia?


- High Performance
- Compiles to Native Code
- Dynamically Typed
- High Level Expressibility
- Parallel Computing Support



[1]



[2]



[3]

Supported Parallelism

- Julia Tasks (coroutines)
- Multi-Threading
- Multi-Core / Distributed Processing



PARALLEL COMPUTING IN JULIA

Julia Tasks

Task Lifecycle

Key Features

- Suspend and Resume
- Query Task State
- Wait on Completion
- **But**, technically not parallelisation as tasks are multiplexed in single OS thread

Key Functions

- `schedule(task)` #schedule a task
- `sleep(seconds)` #sleep current task
- `notify(condition, val, all, error)` #wake up tasks waiting on condition passing val
- `yield(task, args)` #immediately schedules task (then calls scheduler)

```
# verbose method
julia> example() = print("Julia Tasks!")

julia> mytask = Task(example)

julia> istaskscheduled(mytask)
false

julia> schedule(mytask)
Julia Tasks!

julia> istaskdone(mytask)
true

# shorthand method
julia> mytask = @task print("Julia Tasks!")

julia> schedule(mytask)
Julia Tasks!
```

Julia Tasks

Inter-Task Communication

Channels

- Waitable FIFO queue
- Multiple readers and writers concurrently
- `put!` blocks when channel is full
- `take!` blocks when channel is empty
- Maximum size argument `Channel(size)`

Key Functions

- `put!(channel, value)` #adds to queue
- `take!(channel)` #get and remove first item from queue (equivalent to `pop()` in Java)
- `fetch(channel)` #get first item from queue (equivalent to `peek()` in Java)
- `bind(channel, task)` #link lifecycle of channel with task

```
c1 = Channel{32} # input data
c2 = Channel{32} # output data

[...] # prepare input data

function foo()
    while true
        data = take!(c1)
        [...] # process data
        put!(c2, result) # write out result
    end
end
```

Multi-Threading

Multi-Threading Functions

Overview

- Shared memory parallelism
- Allows compute bound tasks to execute in parallel
- Execute for loop iterations simultaneously
- Uses block allocation of iteration space
- **But**, the number of julia threads must be defined prior to start up using `'export JULIA_NUM_THREADS=4'`

Key Functions

- `@threads` #indicate multi-threaded region
- `nthreads()` #number of threads available
- `threadid()` #id if executing thread

```
julia> using Base.Threads

julia> nthreads()
4

julia> a = zeros(10)
#output omitted

julia> @threads for i = 1:10
                a[i] = threadid()
            end

julia> a
10-element Array{Float64,1}:
 1.0
 1.0
 1.0
 2.0
 2.0
 2.0
 3.0
 3.0
 4.0
 4.0
```

Multi-Threading

Atomic Operations

Overview

- Thread-safe way to avoid race conditions
- Wraps values and provides methods
- Create using `Atomic{T} (val)`
- Wrapped value **must** be a primitive type

Key Functions

- `atomic_add!(x, val)` #add val to x
- `atomic_sub!(x, val)` #subtract val from x
- `atomic_and!(x, val)` #bitwise and of x and val
- `atomic_or!(x, val)` #bitwise or of x and val
- `atomic_cas!(x, old, new)` #compare and set x
- `atomic_xchg!(x, new)` #exchange value in x

```
julia> using Base.Threads
```

```
julia> x = Atomic{Int}(2)  
Atomic{Int64}(2)
```

```
julia> atomic_add!(x, 10)  
2
```

```
julia> x  
Atomic{Int64}(12)
```


Multi-Core / Distributed Processing

Workers and Local vs Remote

Workers

- All processes for parallel execution (except process 1) are referred to workers
- We define the number of workers on startup using ``./julia -p n`` where `n` is #workers
- Workers can be added / removed on the fly using `addprocs(num)` and `rmprocs(num)`

Paradigm

- Modified version of the message passing paradigm
- One side message passing - programmer explicitly manages one of the two processors

Local vs Remote

- **Remote References** are references which refer to an object located on another processor.
- **Remote Calls** are requests by a processor to call some function on another processor. This returns a remote reference (Future) which can be used to `wait()` on, or `fetch()` the result from.
- Remote calls return immediately and the calling processor proceeds to its next operation. However, if you need to wait on the result, you can use `wait()` to block until completion of the remote call.

Multi-Core / Distributed Processing

Code Availability, Spawning, Sync and Async

Code Availability

- Code must be available on any processor that runs it
- We use `@everywhere` in the include call to make it available to all processors
- This does **not** bring it into scope

```
julia> @everywhere include("Test.jl")
loaded
      From worker 3:    loaded
      From worker 2:    loaded

julia> test = @spawn myid()
Future(2, 1, 5, nothing)

julia> test1 = @spawnat 2 myid()
Future(2, 1, 3, nothing)
```

Spawning

- `@spawn expr # run exp` expression on an automatically chosen processor
- `@spawnat proc expr # run` expression on `proc` processor

Sync & Async

- `@async task # add` task to local machine scheduler
- `@sync task # blocks` until completion of all `@async`, `@spawn`, `@spawnat`, `@distributed`

Multi-Core / Distributed Processing

Parallel Maps and Loops

Parallelising Loops

- Loops can be parallelised on any number of processors with the `@distributed` macro
- This can optionally carry out a reduction based on the specified reduction. i.e. `@distributed (reduction)`

```
# count num heads
num = @distributed (+) for i = 1:20000
    Int(rand{Bool})
end
```

Shared Arrays

- Parallelising loops does not work as intended when you need to work with shared arrays since each processor gets its own copy
- To solve this, we use a `SharedArray` which maps into shared memory
- `SharedArray{T,N}` # array of type T and size N

```
using SharedArrays

a = SharedArray{Float64}(10)
@distributed for i = 1:10
    a[i] = i
end
```



PARALLELISATION OF GRAPH ALGORITHMS


Parallelising Graph Algorithms

Graph applications

- Social networks
- Road networks
- Communication networks

Why parallelise?

- Very useful algorithms
- Commonly run with huge amounts of data
- Gain **faster runtimes**



A **graph** is a set of nodes and a set of edges

Parallelising Graph Algorithms

Common graph algorithms

- Dijkstra - shortest path
- Breadth-first search (BFS) - graph traversal
- Depth-first search (DFS) - graph traversal
- Prim's algorithm - minimum spanning tree
- Many more

Selected for parallelisation

- Prim's algorithm
- Breadth-first search

Prim's Algorithm - Sequential

Given a weighted graph, find its minimum spanning tree

```
procedure PRIM_MST( $V, E, w, r$ )
begin
   $V_T := \{r\};$ 
   $d[r] := 0;$ 
  for all  $v \in (V - V_T)$  do
    if edge  $(r, v)$  exists set  $d[v] := w(r, v);$ 
    else set  $d[v] := \infty;$ 
  while  $V_T \neq V$  do
    begin
      find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$ 
       $V_T := V_T \cup \{u\};$ 
      for all  $v \in (V - V_T)$  do
         $d[v] := \min\{d[v], w(u, v)\};$ 
      endwhile
    endwhile
  end PRIM_MST
```

Algorithm Features:

- **Builds MST incrementally**
Adds one new node each iteration..
- **Distance vector**
 $d[v]$ represents edge with smallest weight of those ending at node v and originating from a node in MST.
- **Greedy approach**
At each iteration adds the minimal edge to the MST.

Prim's Algorithm - Sequential

Given a weighted graph, find its minimum spanning tree

```
procedure PRIM_MST( $V, E, w, r$ )  
begin  
   $V_T := \{r\};$   
   $d[r] := 0;$   
  for all  $v \in (V - V_T)$  do  
    if edge  $(r, v)$  exists set  $d[v] := w(r, v);$   
    else set  $d[v] := \infty;$   
  while  $V_T \neq V$  do  
    begin  
      find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$   
       $V_T := V_T \cup \{u\};$   
      for all  $v \in (V - V_T)$  do  
         $d[v] := \min\{d[v], w(u, v)\};$   
    endwhile  
  end PRIM_MST
```


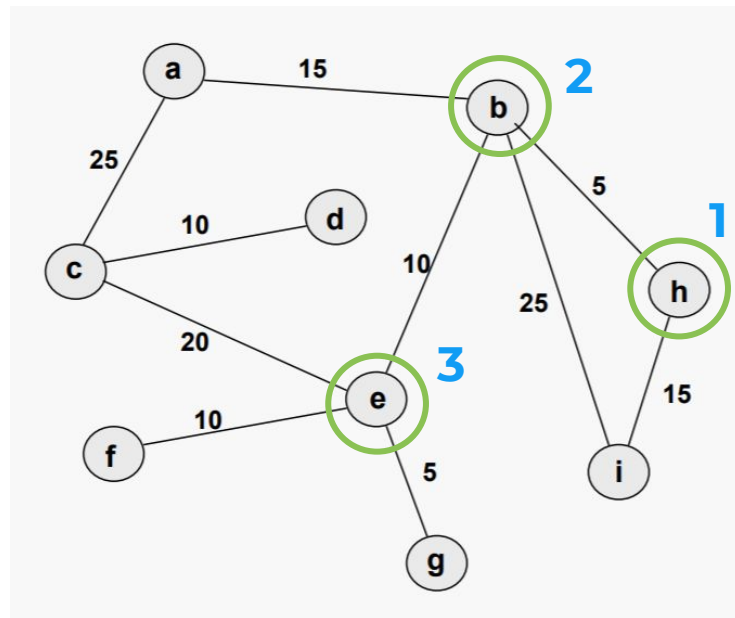


Image retrieved from: <http://courses.cs.vt.edu/~cs3114/Fall10/Notes/T22.WeightedGraphs.pdf> 16

Prim's Algorithm - Sequential

Given a weighted graph, find its minimum spanning tree

```
procedure PRIM_MST( $V, E, w, r$ )  
begin  
   $V_T := \{r\};$   
   $d[r] := 0;$   
  for all  $v \in (V - V_T)$  do  
    if edge  $(r, v)$  exists set  $d[v] := w(r, v);$   
    else set  $d[v] := \infty;$   
  while  $V_T \neq V$  do  
    begin  
      find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$   
       $V_T := V_T \cup \{u\};$   
      for all  $v \in (V - V_T)$  do  
         $d[v] := \min\{d[v], w(u, v)\};$   
      endwhile  
    end PRIM_MST
```



Prim's Algorithm - Parallel

Given a weighted graph, find its minimum spanning tree

```
procedure PRIM_MST( $V, E, w, r$ )
```

```
begin
```

```
   $V_T := \{r\};$ 
```

```
   $d[r] := 0;$ 
```

```
  for all  $v \in (V - V_T)$  do
```

```
    if edge  $(r, v)$  exists set  $d[v] := w(r, v);$ 
```

```
    else set  $d[v] := \infty;$ 
```

```
  while  $V_T \neq V$  do
```

```
    begin
```

```
      find a vertex  $u$  such that  $d[u] := \min\{d[v] | v \in (V - V_T)\};$ 
```

```
       $V_T := V_T \cup \{u\};$ 
```

```
      for all  $v \in (V - V_T)$  do
```

```
         $d[v] := \min\{d[v], w(u, v)\};$ 
```

```
    endwhile
```

```
end PRIM_MST
```

Outer loop difficult to parallelise.

Inner steps are parallelisable


Prim's Algorithm - Parallel

Data partitioning

- Assumes adjacency matrix graph representation
- Split up matrix and distance vector into chunks
- Each processor handles its own distance vector and adjacency matrix partition.

Inner loop steps

- Each processor calculates local minimum edge based on their respective distance vectors.
- Local minima are “reduced” to a global minimum
- The global minimum is added to the MST.
- Update distance vector by looping over the row of the added node, parallelising on columns.



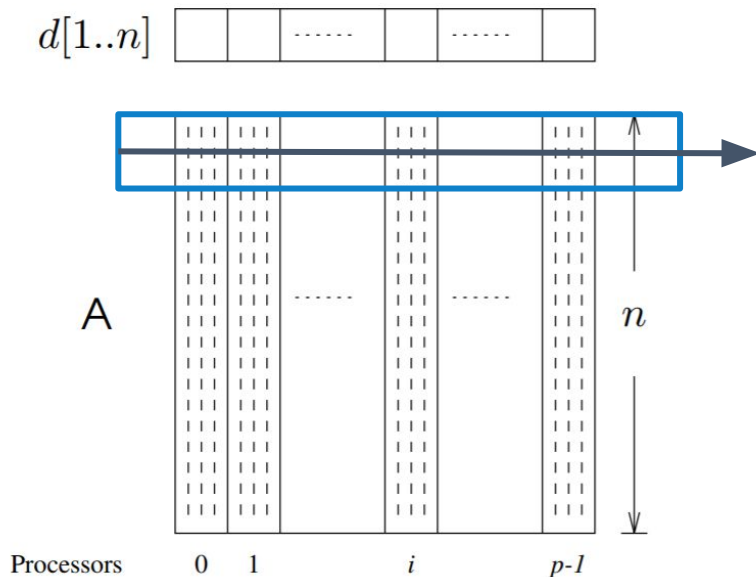
Prim's Algorithm - Parallel

Data partitioning

- Assumes adjacency matrix graph representation
- Split up matrix and distance vector into chunks
- Each processor maintains its own distance vector partition.

Inner loop steps

- Each processor calculates local minimum edge based on their respective distance vectors.
- Local minima are “reduced” to a global minimum
- The global minimum is added to the MST.
- Update distance vector by looping over the row of the added node, parallelising on columns.



Breadth-First Search - Sequential

Given a graph and a start node, visit every node reachable from the start node

High-level Approach

Visit each node **layer by layer**:

- Visit the source node to begin
- As we visit a new node, mark all its neighbours as “to visit”.
- Visit all nodes marked “to visit” until there are none remaining

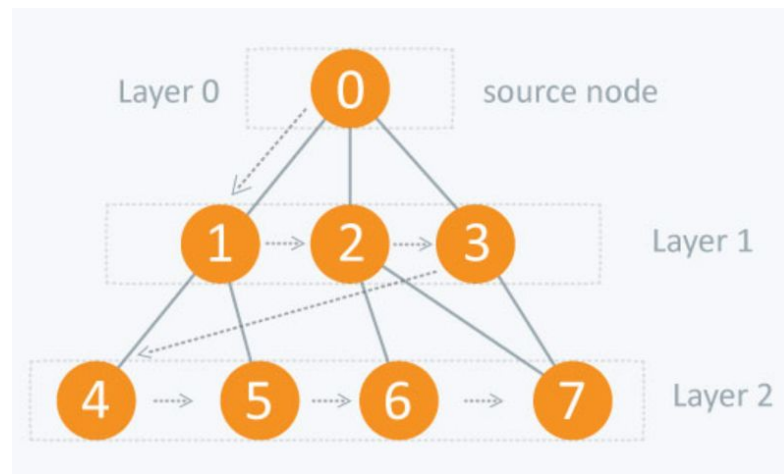


Image retrieved from: <https://www.hackerearth.com/practice/algorithms/graphs/breadth-first-search/tutorial/>²¹

Breadth-First Search - Sequential

Given a graph and a start node, visit every node reachable from the start node

Implementation

- Boolean array to track visited / discovered nodes
- Collections for the current level, and the next level.
- Alternatively use a single queue - harder to parallelise

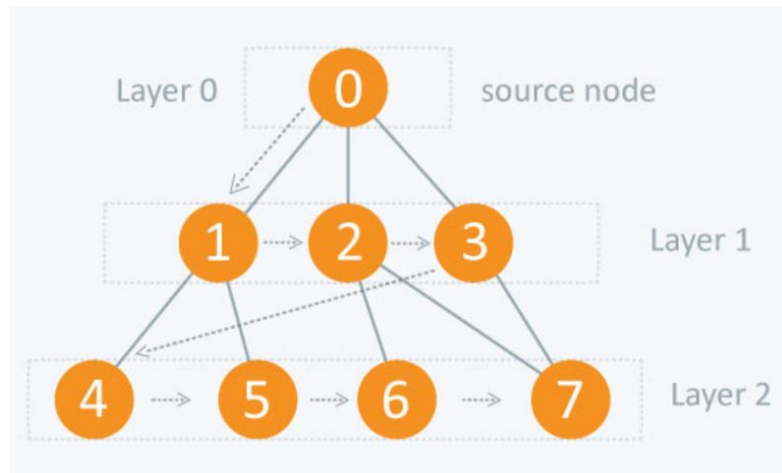



Image retrieved from: <https://www.hackerearth.com/practice/algorithms/graphs/breadth-first-search/tutorial/>

Breadth-First Search - Parallel

Approach

- Like prim's, can't parallelise the “outer loop”
- Parallelise on each level:
 - Split level into p equal-sized chunks
 - Each processor visits each node in its own chunk
 - Possible contention adding to next level collection and checking visited array.



Both P1 and P2:

- If neighbour is not discovered
- Set to discovered
- Add to next level



PARALLEL GRAPH IMPLEMENTATION & BENCHMARKING

Benchmarking in Julia

BenchmarkTools

- `@benchmark`
- `@btime`
- `@belapsed`

`@benchmark`

- Specify parameters
- Returns a `Trial` object

```
julia> @benchmark rand() samples=1000
BenchmarkTools.Trial:
  memory estimate:  0 bytes
  allocs estimate:  0
  -----
  minimum time:      4.908 ns (0.00% GC)
  median time:       5.664 ns (0.00% GC)
  mean time:         6.545 ns (0.00% GC)
  maximum time:      156.695 ns (0.00% GC)
  -----
  samples:            1000
  evals/sample:       1000
```

Benchmarking in Julia

Statistics

- `minimum(Array arr)`
- `median(Array arr)`
- `mean(Array arr)`
- `maximum(Array arr)`

```
julia> benchmark = @benchmark rand() samples=1000
...

julia> dump(benchmark)
BenchmarkTools.Trial
  params: BenchmarkTools.Parameters
    seconds: Float64 5.0
    samples: Int64 1000
    evals: Int64 1000
    overhead: Float64 0.0
    gctrail: Bool true
    gcsample: Bool false
    time_tolerance: Float64 0.05
    memory_tolerance: Float64 0.01
  times: Array{Float64}((1000,)) [4.909 ... 30.206]
  gctimes: Array{Float64}((1000,)) [0.0 ... 0.0]
  memory: Int64 0
  allocs: Int64 0

julia> println(median(benchmark))
TrialEstimate(6.041 ns)
```

Benchmarking in Julia

Setup

- 1000 samples
- 120 second timeout
- A set of randomly generated graph of various sizes

CPU

- Intel i5-4570 @ 3.20GHz
- Can be pushed to 3.60GHz
- 4 physical cores and threads

Benchmarking in Julia

Nested Loop

```
cheapest, index

for each node already visited
  for every other node
    if edge is smaller than cheapest
      update cheapest and index
    end if
  end for
end for
```

Sequential

- 1000 node graph
- 625 seconds

Parallel

- 1000 node graph
- 25 seconds

Benchmarking in Julia

Parallel Distance Vector Implementation

cheapestNode(distanceVector, mst)

```
addprocs(n)

...

cheapest = @distributed min for i = 1:length(distanceVector)
    if node not in mst
        ...
        calculate cheapest node
        ...
    end
    (minCost, nodeIndex)
end

return cheapest[2]
```

Benchmarking in Julia

Parallel Distance Vector Implementation


updateVector(...)

```
addprocs(n)
```

```
...
```


```
@sync @distributed for i = 1:length(distanceVector)
    if new node edge cost cheaper
        update distance
    end
end
end
```

Preliminary Results



Sequential

Distance Vector with set
of traversed nodes



Parallel

Parallelised version using
traversed nodes

Benchmarking in Julia

Parallel Distance Vector (with set of nodes) Implementation

cheapestNode(distanceVector, remainingNodes)

```
addprocs(n)

...

cheapest = @distributed min for i in collect(remainingNodes)
    (distanceVector[i], i)
end


return cheapest[2]
```


Preliminary Results



Sequential v2

Distance Vector with set
of remaining nodes




Parallel v2

Parallelised version using
remaining nodes

Preliminary Results

Cheapest Node Function



Future Work

- Further investigate and optimise parallelisation methods
- Implement BFS algorithms
- Investigate more complex graph algorithms
- Compare performance against other languages

THANK YOU

Questions?

REFERENCES

- [1] Julia. (2016). *Machine learning* [Screenshot]. Retrieved from <https://julialang.org/>
- [2] Julia. (2016). *Data science* [Screenshot]. Retrieved from <https://julialang.org/>
- [3] Julia. (2016). *Scientific Computing* [Screenshot]. Retrieved from <https://julialang.org/>